

# DHEERAJ KUMAR

Tempe, AZ | (480) 931-8330 | [dkumar70@asu.edu](mailto:dkumar70@asu.edu) | [linkedin.com/in/dheeraj-k-232a6a235](https://linkedin.com/in/dheeraj-k-232a6a235) | [github.com/dheeraj7000](https://github.com/dheeraj7000)

## EDUCATION

Arizona State University — M.S. in Computer Science

Tempe, AZ | Expected Dec 2026

GPA: 3.56 | Coursework: Knowledge Representation, Topics in NLP, Statistical Machine Learning, Data Processing at Scale

## TECHNICAL SKILLS

**Languages:** Python, Rust, SQL, JavaScript/TypeScript

**AI / ML:** PyTorch, PyTorch Geometric, Graph Neural Networks, Transformers, Model Fine-Tuning, Embeddings, NLP, Statistical Machine Learning

**LLM & Agentic Systems:** Agentic System Design, Multi-Agent Orchestration, RAG, Tool Calling, MCP (Model Context Protocol), Prompt Engineering, LLM Evaluation, Agent Memory Architectures, Prompt Injection Defense, AI Threat Modeling (OWASP LLM Top 10, MITRE ATLAS)

**Engineering:** FastAPI, REST APIs, JSON-RPC, Redis, Docker, Git, Linux, Playwright, Static Analysis

## EXPERIENCE

Pinpoint Technologies LLC — ML Research / Engineering

Remote | May 2026 – Present

- Designed and implemented graph neural network models with PyTorch Geometric for relational network data, improving prediction accuracy over non-graph baselines by enabling richer relational representations
- Developed components of proprietary Rust-based security analysis tooling and blockchain data pipelines (project specifics under NDA)
- Built data preprocessing and feature engineering pipelines for structured and unstructured datasets, enabling reproducible, version-controlled evaluation runs

Time Pilot — NLP Engineer Intern

Remote | Sep 2024 – Mar 2025

- Developed and integrated RESTful APIs with FastAPI connecting ML models to front-end applications, improving system integration and scalability for end users
- Fine-tuned a custom multitask transformer model serving multiple classification and text-generation tasks from a single deployed model, reducing serving overhead versus per-task models

## PROJECTS

AnyAPI — Self-Hosted Multi-Provider LLM API Server | Python, Playwright, JSON-RPC

[dheeraj7000.github.io/anyapi](https://dheeraj7000.github.io/anyapi)

- Built an open-source, pip-installable API server that exposes nine free LLM web providers (DeepSeek, ChatGPT, Gemini, Perplexity, Copilot and others) behind a single streaming JSON-RPC interface over a Unix domain socket — no API keys required
- Engineered a long-lived Playwright/Chromium daemon with per-provider rate limiting, hot cookie/auth reload, health checks, and conversation persistence, letting many lightweight clients (CLI, scripts, bots, agent pipelines) share one browser session
- Shipped a provider-plugin architecture, token-level streaming event protocol, CLI reference client, and public documentation site to support external contributors

CRM AI Agent — LLM Security Architecture Review | LLM Security, Threat Modeling

- Performed end-to-end security review of a CRM-integrated LLM agent architecture, identifying prompt injection gaps, trust-boundary violations, token-caching inversion, and a broken SSE stream-resume mechanism
- Mapped findings to OWASP LLM Top 10 and MITRE ATLAS; delivered a prioritized remediation report covering tool-call authorization and data-exfiltration paths

## OPEN SOURCE CONTRIBUTIONS

SkillNet — Agent Skill Infrastructure (ZJUNLP / Zhejiang University NLP Lab) | [Merged PR #33](#)

Aug 2026

- Remediated a prompt-injection flaw in the LLM-as-judge skill evaluation gate, where untrusted third-party skill content was interpolated into the judge prompt with no trust boundary, letting a skill influence its own safety rating
- Contributed a dependency-free static pre-screen module (11 detection rules), nonce-fenced prompt sections, and 26 tests including false-positive guards for security-domain skills

## TECHNICAL REPORTS & PREPRINTS

- Carbowitz, R., Kumar, D. "SPOQ: Specialist Orchestrated Queuing for Multi-Agent Software Engineering." [arXiv:2606.03115](https://arxiv.org/abs/2606.03115) (2026)
- Kumar, D. "GCAF-Net: Gated Cross-Attention Fusion for Efficient Robotic Vision-Language Reasoning." [ResearchGate](#)
- Kumar, D., Rana, V. "Ontology Engineering for Trustworthy AI: A Production-Grounded Survey of Knowledge Graphs in the LLM Era." [ResearchGate](#)

## ACHIEVEMENTS

Prime Minister Scholarship (full tuition waiver) · 1st Place, AI Scholar Principled AI Hackathon · 1st Place, Dutch Bros × ASU AI Broista Challenge 2025